

Class06: R Functions

Hannah Kim

4/21/23

In this class we will develop our own R function to calculate average grades in a fictional class.

We will start with a simplified version of the problem, just calculating the average grade of one student

Simplified Version

```
# Example input vectors to start with

student1 <- c(100, 100, 100, 100, 100, 100, 100, 90)
student2 <- c(100, NA, 90, 90, 90, 90, 97, 80)
student3 <- c(90, NA, NA, NA, NA, NA, NA, NA)
```

We are going to start by calculating the average score of the homeworks.

```
mean(student1)
```

```
[1] 98.75
```

To get the minimum score we can use which.min.

```
student1
```

```
[1] 100 100 100 100 100 100 100 90
```

```
which.min(student1)
```

```
[1] 8
```

I can do the average of the first 7 homework scores:

```
mean(student1[1:7])
```

```
[1] 100
```

Another way to select the first 7 homework scores:

```
student1[1:7]
```

```
[1] 100 100 100 100 100 100 100
```

```
student1[-8]
```

```
[1] 100 100 100 100 100 100 100
```

Another way to drop the lowest score:

```
#goes through student1 scores and removes the minimum value  
student1_drop_lowest <- student1[-which.min(student1)]
```

I can get the mean of the homework scores after dropping the lowest score by doing.

```
mean(student1_drop_lowest)
```

```
[1] 100
```

We have our first working snippet of code!

Let's try to generalize it to student2:

```
student2 <- c(100, NA, 90, 90, 90, 90, 97, 80)  
student2_drop_lowest<-student2[-which.min(student2)]  
student2_drop_lowest
```

```
[1] 100 NA 90 90 90 90 97
```

There is a way to calculate the mean dropping missing values

```
student3 <- c(90, NA, NA, NA, NA, NA, NA, NA)
mean(student3, na.rm = TRUE)
```

```
[1] 90
```

We want to know the position of the NAs. So, for student2 we can use the following.

```
student2 <- c(100, NA, 90, 90, 90, 90, 97, 80)
which(is.na(student2))
```

```
[1] 2
```

For student 3:

```
which(is.na(student3))
```

```
[1] 2 3 4 5 6 7 8
```

For student 2:

```
student2
```

```
[1] 100 NA 90 90 90 90 97 80
```

```
which(is.na(student2))
```

```
[1] 2
```

```
student2[ is.na(student2) ] <- 0
student2
```

```
[1] 100 0 90 90 90 90 97 80
```

If I use the same for student 3

```
student3[ is.na(student3) ] <- 0
student3
```

```
[1] 90  0  0  0  0  0  0  0
```

```
mean(student3)
```

```
[1] 11.25
```

This is going to be our final working snippet of code for all students (with and without NA values)

```
student3 <- c(90, NA, NA, NA, NA, NA, NA, NA)
student3[is.na(student3)] <- 0
student3_drop_lowest <- student3[-which.min(student3)]
mean(student3_drop_lowest)
```

```
[1] 12.85714
```

Q1

We can write it as a function:

```
#creating a function that requires an array x
grade <- function(x)
{
  #finds the index of the value that is NA and then changes it to 0
  x[is.na(x)] <- 0
  #creates a variable that stores the average of student scores without the lowest score
  x_drop_lowest <- x[-which.min(x)]
  mean(x_drop_lowest) }
```

Let's apply the function

```
grade(student1)
```

```
[1] 100
```

```
grade(student2)
```

```
[1] 91
```

```
grade(student3)
```

```
[1] 12.85714
```

Let's apply our function to a gradebook from this URL: "https://tinyurl.com/gradeinput"

```
URL <- "https://tinyurl.com/gradeinput"
gradebook <- read.csv(URL, row.names = 1)
head(gradebook)
```

	hw1	hw2	hw3	hw4	hw5
student-1	100	73	100	88	79
student-2	85	64	78	89	78
student-3	83	69	77	100	77
student-4	88	NA	73	100	76
student-5	88	100	75	86	79
student-6	89	78	100	89	77

Let's apply my function grade to the gradebook using apply and running it by rows using MARGIN = 1.

```
apply(gradebook, 1, grade)
```

student-1	student-2	student-3	student-4	student-5	student-6	student-7
91.75	82.50	84.25	84.25	88.25	89.00	94.00
student-8	student-9	student-10	student-11	student-12	student-13	student-14
93.75	87.75	79.00	86.00	91.75	92.25	87.75
student-15	student-16	student-17	student-18	student-19	student-20	
78.75	89.50	88.00	94.50	82.75	82.75	

Q2

We can write it as a function

```
max(apply(gradebook,1,grade))
```

```
[1] 94.5
```

The maximum score is 94.5

```
which.max(apply(gradebook,1,grade))
```

```
student-18  
18
```

The student getting the maximum overall score was student 18.

Q3

First we are going to mask NA values with zeros

```
#replacing all NAs with 0 but applying it to dataframe instead of array  
gradebook[is.na(gradebook)] <- 0
```

Now we apply the mean function to the gradebook

```
apply(gradebook,2,mean)
```

```
hw1 hw2 hw3 hw4 hw5  
89.00 72.80 80.80 85.15 79.25
```

The toughest homework will be homework 2 considering the mean and considering the missing homework as 0.

Having zeros for missing homework is too strict and is not a good representation of the homework difficulty.

One thing we can do is remove the missing values.

```
gradebook<- read.csv(URL, row.names = 1)  
apply(gradebook,2,mean,na.rm = TRUE)
```

```
hw1 hw2 hw3 hw4 hw5  
89.00000 80.88889 80.80000 89.63158 83.42105
```

Q4. From your analysis of the gradebook, which homework was most predictive of overall score


(i.e. highest correlation with average grade score)?

If we use the median instead of the mean as a measure of overall score:

```
apply(gradebook, 2, median, na.rm = TRUE)
```

hw1	hw2	hw3	hw4	hw5
89.0	72.5	76.5	88.0	78.0

```
boxplot(gradebook)
```



```
overall_grades <- apply(gradebook, 1, grade)
overall_grades
```

student-1	student-2	student-3	student-4	student-5	student-6	student-7
91.75	82.50	84.25	84.25	88.25	89.00	94.00

student-8	student-9	student-10	student-11	student-12	student-13	student-14
93.75	87.75	79.00	86.00	91.75	92.25	87.75
student-15	student-16	student-17	student-18	student-19	student-20	
78.75	89.50	88.00	94.50	82.75	82.75	

```
cor(gradebook$hw1, overall_grades)
```

```
[1] 0.4250204
```

```
apply(gradebook,2,cor, y = overall_grades)
```

hw1	hw2	hw3	hw4	hw5
0.4250204	NA	0.3042561	NA	NA

```
apply(gradebook,2,cor, y = overall_grades)
```

hw1	hw2	hw3	hw4	hw5
0.4250204	NA	0.3042561	NA	NA